

Introduction to Large Language Models (LLM)

Large language models are deep neural networks trained on vast text data.

They can generate human-like text, answer questions, and solve complex problems.

However, LLMs suffer from hallucination and lack private domain knowledge.

To overcome this limitation, Retrieval-Augmented Generation (RAG) is introduced.

Chapter 2: Retrieval-Augmented Generation (RAG)

A typical RAG pipeline contains three major stages:

1. Indexing: Documents are chunked and transformed into vector embeddings.
2. Retrieval: When a user query arrives, relevant chunks are retrieved.
3. Generation: Chunks are provided to the LLM as ground truth context.